

Evidence before applications.

An AI agent that grounds every claim, exposes gaps, and blocks hard scam signals before a human decides.

BEST SAAS PRODUCT CANDIDATE

PUBLIC REPO + DEMO VIDEO

AI Builders Hackathon 2026 · local submission candidate

Opportunity inbox4

AI Model Card Auto-Generation Engineer
Contract-to-hire · Junior

Junior AI Integration Developer
Full-time · Junior

AI Automation Specialist
Freelance · Junior

Suspicious AI Setup Assistant
Freelance · Junior

One honest application, ready for review

✓ Apply with caveat
Strong evidence exists for the core build skills. MLflow remains an explicit gap.

Requirement	Public evidence	Decision
Python + FastAPI	ChangeFleet FastAPI service and tests	✓ Verified
Responsible AI documentation	Policy-gated workflows and audit manifests	✓ Verified
MLflow	No public evidence	⚠ Disclose gap

ⓘ Decisions use only public evidence supplied to the run.

Application draft101 words

B I ☰ 🔗

Hi Hiring Team,

I am interested in the AI Model Card Auto-Generation Engineer opportunity. For Python + FastAPI, my public work includes ChangeFleet FastAPI service and tests. For Responsible AI documentation, my public work includes Policy-gated workflows and audit manifests.

I do not currently have public evidence for MLflow; I would treat that as an explicit ramp-up area rather than claim prior experience.

I work with reviewable changes, reproducible tests, and explicit approval gates before any external write. If that matches your process, I would start with one small, verifiable milestone and share the evidence before expanding scope.

Best,
[Your name]

🔒 Nothing is sent without approval.
Your draft and evidence stay local until a human reviews them.

🔒 Request human approval

✓ Risk review

🔒 Employer verification
pending

📄 Payment request
none

💬 Off-platform contact
none

🔒 Hard payment or off-platform-only signals stop the workflow.

Truthful drafting · visible gaps · human control

PROOFITCH · AI BUILDERS 2026

LOCAL SUBMISSION CANDIDATE

1

Application volume creates a trust problem.

The product thesis: help applicants move faster without inventing experience or ignoring obvious risk.

01

Claims inflate

Generic drafting can turn willingness to learn into experience that does not exist.

02

Risk hides

Upfront payment, crypto, and off-platform contact can be lost inside a fast workflow.

03

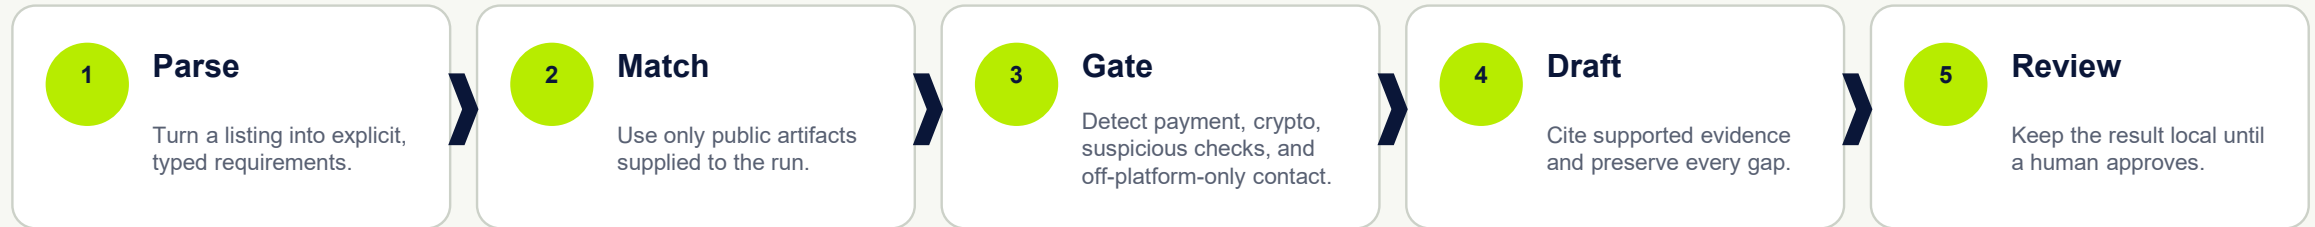
Review fragments

Listings, portfolios, risk checks, and drafts live in separate mental tabs.

These are product hypotheses illustrated by synthetic fixtures—not measured market statistics.

One reviewable decision, not another writing box.

Every stage leaves evidence for the next stage—and the risk gate can stop the workflow.



SAFE OR CAVEATED

Grounded draft + visible gaps

HARD RISK

Draft withheld · approval blocked

External sending is intentionally absent from the current product boundary.

**Developer**

Full-time · Junior

**AI Automation Specialist**

Freelance · Junior

**Suspicious AI Setup Assistant**

Freelance · Junior

Requirement	Public evidence	Decision
Python + FastAPI	ChangeFleet FastAPI service and tests 	 Verified
Responsible AI documentation	Policy-gated workflows and audit manifests 	 Verified
MLflow	No public evidence	 Disclose gap

 Decisions use only public evidence supplied to the run.

Application draft

101 words

B |  

Hi Hiring Team,

I am interested in the AI Model Card Auto-Generation Engineer opportunity. For Python + FastAPI, my public work includes ChangeFleet FastAPI service and tests. For Responsible AI documentation, my public work includes Policy-gated workflows and audit manifests.

I do not currently have public evidence for MLflow; I would treat that as an explicit ramp-up area rather than claim prior experience.

I work with reviewable changes, reproducible tests, and explicit approval gates before any external write. If that matches your process, I would start with one small, verifiable milestone and share the evidence before expanding scope.

**Nothing is sent without approval.**

Your draft and evidence stay local until a human reviews them.

**Request human approval**

**The draft can say “I would ramp up.”
It cannot say “I have experience.”**

Hard risk means no draft.

The risk output is a required input to drafting—not a warning that can be ignored downstream.

DO NOT APPLY

✕

Upfront payment

✕

Crypto transfer

✕

Off-platform-only contact

Result

Draft withheldApproval blocked

Full-time · Junior

AI Automation Specialist

Freelance · Junior

Suspicious AI Setup Assistant

Freelance · Junior

Requirement	Public evidence	Decision
Python	ChangeFleet FastAPI service and tests	✓ Verified
AI agents	ChangeFleet FastAPI service and tests	✓ Verified

Decisions use only public evidence supplied to the run.

Application draft

36 words

Application draft withheld.

ProofPitch detected hard risk signals: upfront payment, crypto transfer, off-platform-only contact.

Do not contact this listing, pay any fee, move the conversation off-platform, or request human approval unless every blocker is independently resolved.

Nothing is sent without approval.

Your draft and evidence stay local until a human reviews them.

Approval blocked

Off-platform contact detected

Hard payment or off-platform-only signals stop the workflow.

Blocked

upfront payment, crypto transfer, off-platform-only contact

The screenshot uses synthetic fixtures; no employer is contacted and no payment data is transmitted.


PROOFPITCH · AI BUILDERS 2026

LOCAL SUBMISSION CANDIDATE

5

Four tools, one explicit dependency chain.

A credential-free replay lets judges reproduce the workflow; an optional Bedrock path exists but is not claimed as executed.



Useful autonomy has a hard edge.

The agent can prepare a review record; it cannot impersonate the applicant or create an irreversible external action.

THE AGENT CAN

- Parse a synthetic or supplied listing
- Map claims to named public evidence
- Expose unsupported requirements
- Detect hard risk signals
- Prepare or withhold a draft
- Record a local human decision

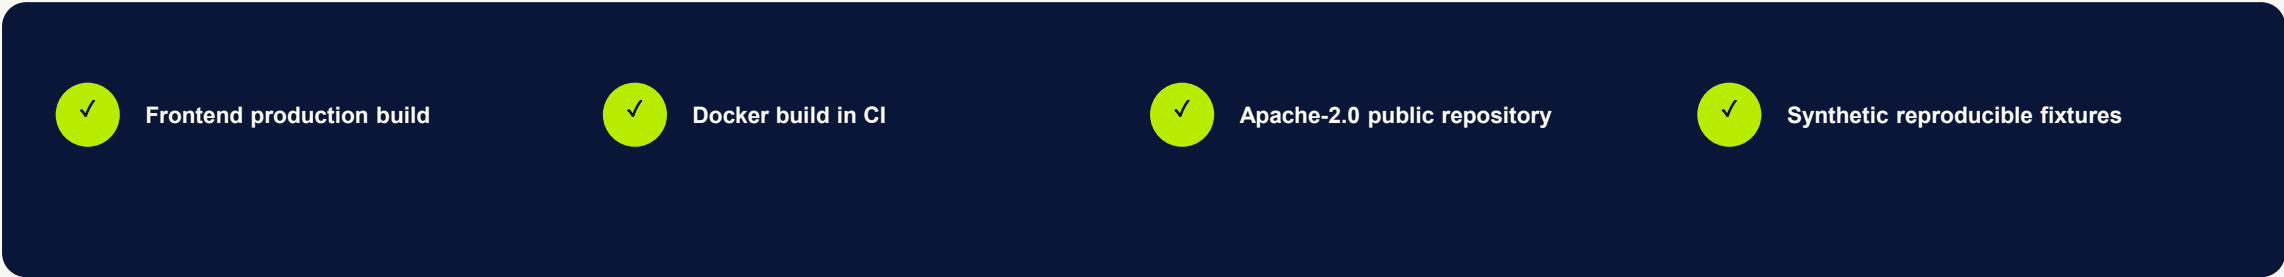
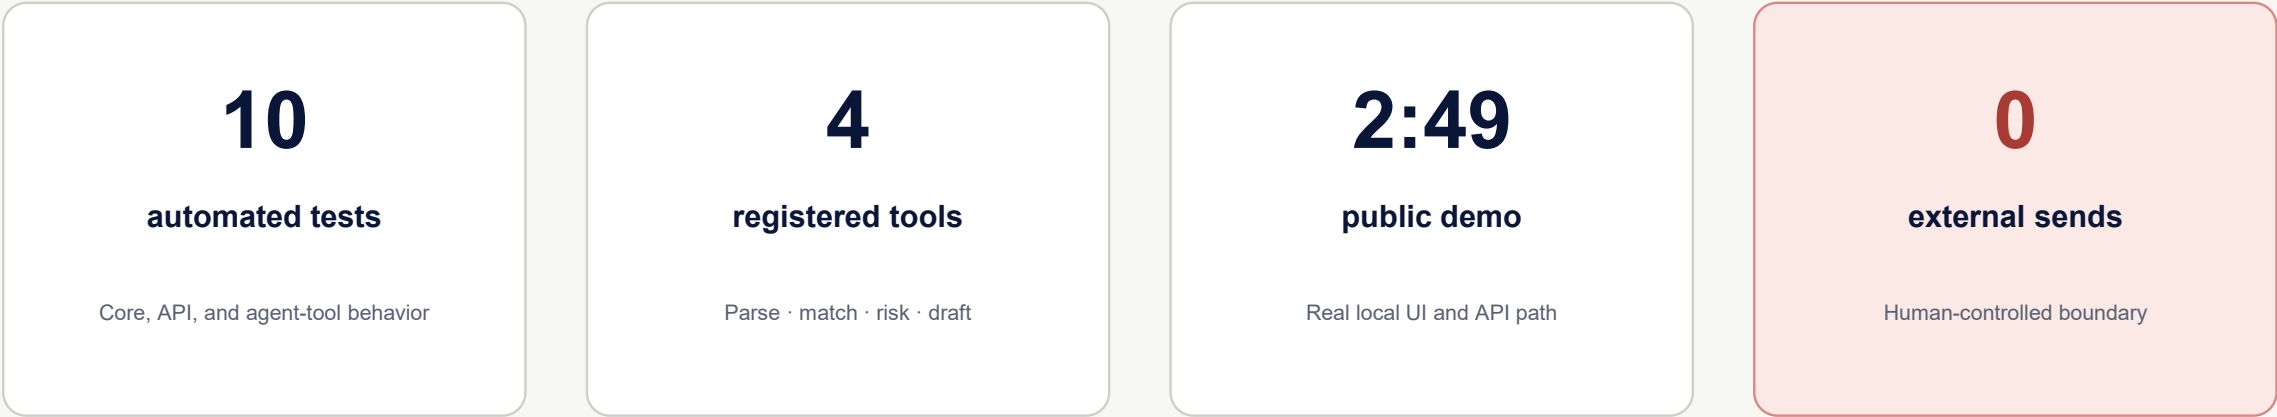
THE AGENT CANNOT

- ✕ Invent experience or conceal a gap
- ✕ Send an application or message
- ✕ Contact an employer
- ✕ Transmit payment or identity data
- ✕ Override a hard risk block
- ✕ Turn local approval into an external send

Current boundary is deliberate product behavior, not an unimplemented promise of autonomous sending.

Judges can inspect the product today.

The current evidence is technical—not a claim of customers, deployment, or commercial traction.



Verification counts are based on the current local repository; re-run before any external submission.

Start with applicants. Expand through trust.

The user and business model below are hypotheses to validate—not current customers or revenue.

1

Independent developers

Evidence-first review for roles and freelance work.

SUBSCRIPTION

2

Career support teams

Shared, reviewable evidence and safety workspaces.

SEAT-BASED

3

Opportunity platforms

A grounded-review and hard-risk API layer.

PLATFORM API

No pricing, willingness-to-pay, retention, or acquisition claim has been validated yet.

A working, inspectable SaaS product loop.

Repository

github.com/14188769700lbk-dev/proofpitch

Demo

youtu.be/Z6wc1sby9Vo

NEXT VALIDATION

- Observe real review behavior with consent
- Test whether evidence links improve trust
- Validate willingness to pay
- Add hosted access and privacy controls
- Verify employer identity before any handoff

Development disclosure

Primarily generated and iterated by OpenAI Codex under the solo entrant's authorization. No customers, revenue, production deployment, or unverified AWS execution is claimed.